

Design and Evaluation of a Framework for Annotating Coloured Petri Net Models with Code Generation Pragmatics

Mikal Hitsøy Henriksen

Master's Thesis

Department of Computer Engineering
Bergen University College
Norway

March 24, 2012
Supervisor
Lars Michael Kristensen

Abstract

Model Driven Development

CPN

Under development: Annotations for code generation

Designing and evaluating framework

Contents

1	Introduction	1
1.1	Model Driven Software Development	1
1.1.1	Code Generation	1
1.2	Related work	1
1.3	Thesis Aim	1
1.4	Thesis organisation	2
1.5	Required Knowledge	2
2	Background	3
2.1	The WebSocket Protocol	3
2.2	Coloured Petri Nets	4
2.2.1	CPN Tools	4
2.3	The WebSocket CPN Model	5
2.3.1	Overview	5
2.3.2	Client Application	8
2.3.3	The Client WebSocket Module	11
2.3.4	Connection	23
2.3.5	The Server WebSocket Layer	24
2.3.6	Server Application	27
3	State Space Analysis of the CPN Web Socket Protocol	29
3.1	State Spaces	29
3.1.1	Strongly Connected Component graph	29
3.1.2	Application of State Spaces	30
3.2	State Space Report	31
3.2.1	Statistics	31
3.2.2	Boundedness Properties	32
3.2.3	Home Properties	33
3.2.4	Liveness Properties	34

3.3	TODO:	34
3.3.1	Error discovery	35
4	Technology and Foundations	37
4.1	Representing CPN Models	37
4.2	Eclipse IDE	38
4.3	Eclipse Modeling Framework (EMF)	40
4.3.1	Graphical Modeling Framework (GMF)	40
4.4	ePNK: Petri Net framework	40
4.5	Access/CPN: Java interface for CPN Tools	41
4.6	Semantic Application Design Language (SADL): Ontologies .	42
4.7	Summary	42
5	Analysis and Design	43
5.1	Requirements	43
5.2	Test cases	43
5.3	Defining Pragmatics	43
5.4	The CPN model type for ePNK	43
5.5	Importing from CPN Tools	44
5.6	Creating annotations	44
5.7	Choosing Pragmatics Sets	44
5.7.1	Creating custom pragmatics	44

Chapter 1

Introduction

Software development paradigms,

1.1 Model Driven Software Development

Short explanation of MDD

1.1.1 Code Generation

Short intro to code generation from models, arguments for

1.2 Related work

References to other works and theses, explaining their approaches

Reference and longer explanation of current work of Simonsen, which defines the domain for the prototype

1.3 Thesis Aim

This thesis seeks to create a framework for annotating CPNs. The framework should be capable of supporting several classes of pragmatics, categorised into General, Domain Specific, and Model Specific Pragmatics.

It should be possible to import CPN models created in CPN Tools. Annotations are placed through a tree editor, and also, if possible, through a graphical diagram editor. Errors highlighting?

Dette bør på sikt formuleres som en hupotese, det skal være et formulert spørsmål som undersøkes.

Finally, there should be an easy way to invoke the code generator from the editor.

1.4 Thesis organisation

The thesis is described below:

Chapter 2: Background Introduction to CPN, showing how to use it by using the WebSocket protocol as an example. Overview of related research. Describe CPN Tools.

Chapter 4: Technology and Foundations Explanation of the Eclipse Platform and its modules, which make up the foundation of our prototype

Chapter 5: Analysis and Design Discussion and detailing of requirements for the editor. Overview and explanation of test cases.

Chapter ??: ?? Describe how the implementation works. Explanation of chosen solutions to requirements.

Chapter ??: ?? Discussion on which requirements have been met. Results from test cases and feedback from users.

Chapter ??: ?? Summary, personal experience, limitations and suggested focus of future work.

1.5 Required Knowledge

The reader is assumed to be familiar with Java programming.

Chapter 2

Background

Bør kanskje gjenvurdere om
det skal hete “Background”

2.1 The WebSocket Protocol

The primary case study for this thesis is the WebSocket protocol [FM11].
From the abstract of the document:

The WebSocket protocol enables two-way communication between a client running untrusted code running in a controlled environment to a remote host that has opted-in to communications from that code.

Fig. 2.1 shows the basic sequence of the WebSocket protocol. To establish a connection, a client sends a specially formatted HTTP request to a server, which replies with a HTTP response. Once the connection is established, the client and server can then freely send WebSocket message frames, until either endpoint sends a control frame with the opcode 0x8 for close and optionally data about the reason for closing. The other endpoint then replies with the same opcode and data, and the connection is considered closed.

From the RFC document:

Conceptually, WebSocket is really just a layer on top of TCP that does the following:

- adds a Web “origin”-based security model for browsers
- adds an addressing and protocol naming mechanism to support multiple services on one port and multiple host names on one IP address;

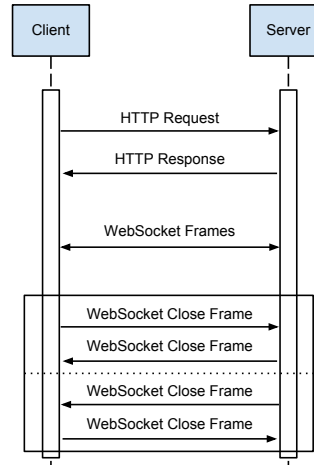


Figure 2.1: Sequence Diagram of the WebSocket protocol

- layers a framing mechanism on top of TCP to get back to the IP packet mechanism that TCP is built on, but without length limits
- includes an additional closing handshake in-band that is designed to work in the presence of proxies and other intermediaries

2.2 Coloured Petri Nets

Common usage: Process and protocol modeling, concurrent programming. Operations: Simulation, verification and analysis. More recently also software design.

The structure of CPN models will be explained gradually through the case study

2.2.1 CPN Tools

Graphical tool used to design CPN models.

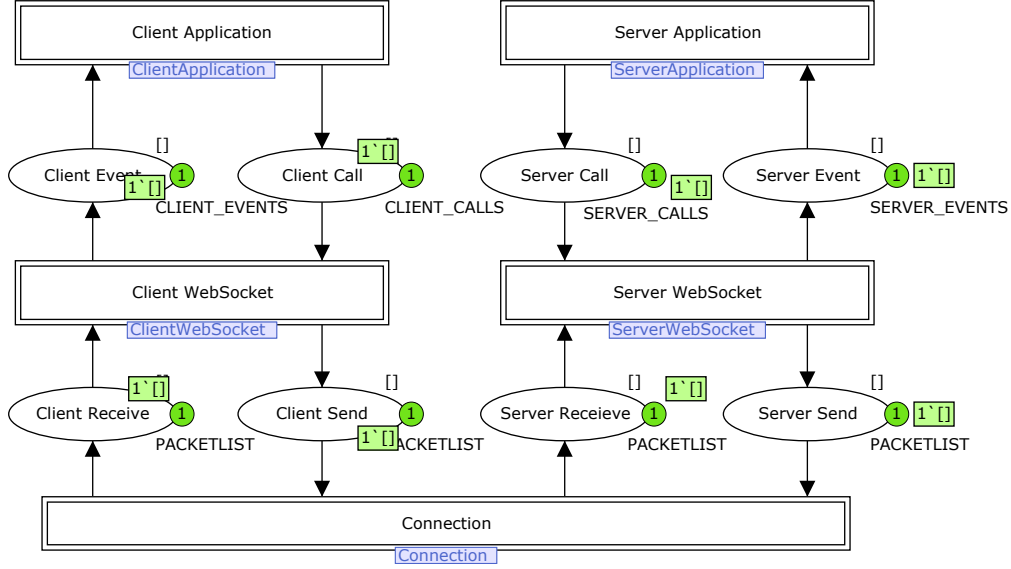


Figure 2.2: Overview of CPN model of the WebSocket protocol

2.3 The WebSocket CPN Model

The websocket protocol is the primary case study of this thesis. In this section the model that has been built is described in full detail.

2.3.1 Overview

Fig. 2.2 shows the top-level Overview module of the WebSocket CPN model.

The model has been laid out to resemble part of the OSI model [Zim80], where Client Application and Server Application each correspond to the top two layers Application and Presentation, Client WebSocket and Server WebSocket correspond to the Session layer, and Connection corresponds to the lower layers Transport, Network, Data Link and Physical.

Model basics

A Petri Net model consist of pages, or modules, that contain Places, Transitions and Arcs.

Places are represented by circles. They can contain colours from a specified colour set, and have markings to define initial colour(s). As an example, in the lower right we have the Server Send place. This place can have

PACKETLIST color tokens, and has an initial marking of an empty list \square . It currently has one colour shown in the green box, $1'\square$, which means one of an empty list. The number in the circle shows the total number of colours in this place.

A colour can mean a single value like the integer 3, the string “Hello” or simple units which resemble Tokens from regular Petri Nets. A colour set is a set of such colours, for example all integers, the integers from 1 to 100, all lowercase letters, or all strings of length 5. These are examples of simple colour sets, as they each consist of basic data types. Compound colour sets are made by combining the simple ones in different ways to create data structures, for example an integer together with a string might describe a postal code and city. This will be described in more detail in the next section.

Transitions are represented by rectangles.

Double-bordered rectangles represent substitution transitions which are bound to sub-modules. These are separate modules that have input and output places that connect to the places on their parent module. This enables us to keep the model structured and readable, and also allows reuse of modules in different parts of the model by instantiating them, which will be described toward the end of the chapter. The details of each submodule is described in the following subsections.

Arcs are directed arrows going between a place and a transition.

The State of the model is defined by the current colors in each place. A Transition can consume colors from places if there is an arc from the place to the transition, and produce new colors in places where there is an arc from the transition to the place. This changes the state of the system.

Declarations

CPN Tools uses the CPN ML language to specify declarations and net inscriptions. This language is an extension of the functional programming language Standard ML, developed at Edinburgh University.

All coloursets, variables, symbolic constants and functions have to be declared before they are used in the model. In CPN Tools these declarations can be grouped together, but they are still parsed sequentially, so if one declaration depends on another, it has to be declared after its dependency.

Colour sets are defined with the following syntax:

```
colset name = <type-specific>;
```

Names are always capitalised in this thesis, but any CPN ML identifier is valid.

The places in the Overview use the following colour set definitions:

Listing 2.1: Overview colour sets

```
colset OPERATION = with CONNECT | TEXT | BINARY | PING | PONG |
    CLOSE;

colset MESSAGE = record Op: OPERATION * Message: STRING;

colset MESSAGES = list MESSAGE;

colset CONN_RESULT = bool with (fail, success);

colset EVENT = union Msg:MESSAGE + ConnRequest + ConnResult:
    CONN_RESULT;

colset EVENTS = list EVENT;

colset CONN_REPLY = bool with (reject, accept);

colset MSG_OR_CONN_REPLY = union Msg':MESSAGE + ConnReply:
    CONN_REPLY;

colset MSGS_OR_CONN_REPLY = list MSG_OR_CONN_REPLY;

colset PACKET = union HttpReq:HTTPREQ + HttpRes:HTTPRES +
    WsFrame:WSFRAME;

colset PACKETLIST = list PACKET;
```

The `OPERATION` colour set is an enumeration that represents the different types of messages that can be passed between the application and the protocol layer. All of these correspond with opcodes used in WebSocket frames, except `CONNECT` which is used for initiating a connection.

A `MESSAGE` is a record that consists of an `OPERATION` and the `STRING` message body. Both data- and control-frames can have messages, although control-frame messages might not be shown to the user. Colours from record color sets can have their component colors referenced by name, in this case `Op` and `Message`. We can also have a list of `MESSAGES` to keep them ordered. Lists will be explained in detail later where they are used and manipulated in the model.

The WebSocket layer also needs a way to notify the application about connection results, which is done with the `EVENT` color set. This is a union color set which can be either a `MESSAGE`, a `ConnRequest` identifier (with no other color associated), or a `CONN_RESULT`. A union color set is used if a place should be able to contain colours from different colour sets, or if such colours should

be handled in the same way at a point in the model. They can also contain simple identifiers (like `ConnRequest`). And like `MESSAGES` we can have a list of `EVENTS`.

The client application sends `MESSAGES` to the WebSocket layer and receives `EVENTS` back. The server application also receives `EVENTS`, but can send either a `MESSAGE` or a `CONN_REPLY`, which is used for connection attempts. The `CONN_REPLY` colourset is simply a boolean value with different names for true and false, for improving readability and semantics. A union of these two types must exist, `MSG_OR_CONN_REPLY`, so the same place can be reused.

TODO: Noe om hvorfor vi bruker kun en plass for dette

Both the client and server WebSocket layers send and receive `PACKET` colors, a union of three types of data. `PACKETS` are abstract and not fully modeled actual network packets, as this is not relevant to how WebSocket works. Their list version `PACKETLIST` These color sets are described in the section for the Connection module.

2.3.2 Client Application

This module (Fig. 2.3) is laid out from top to bottom to loosely show the sequence of operation, as is most of the other modules.


The places at the bottom represent the interface to the WebSocket layer. To simplify the overview model and facilitate easier later modification and expansion, we have only two places that act as input and output. These are called Receive Client Events and Send Client Message, and are tagged In and Out respectively. These are connected to the respective places on the Overview, which are also connected to corresponding places in the WebSocket Library.

Variables

The arcs here have inscriptions on them containing variables. Variables are declared globally and can only be bound to colours of the colourset they are defined for.

When a transition fires, it will bind the tokens it is consuming to the variables of the inscriptions on the arcs from the respective places. These variables can then be used to produce new tokens, by direct manipulation and/or in conditional statements. The expression that will produce the desired result is inscribed on the output arc.

A set of variables has been declared for the simple coloursets as follows:



Listing 2.2: Simple Colourset Variables

```

var u, u1, u2, u3: UNIT;
var b, b1, b2, b3: BOOL;
var i, j, k: INT;
var s, s1, s2, s3: STRING;
var ss, ss1, ss2: STRINGLIST;

```

Several variables for the same coloursets have been created for convenience in the cases where colours of the same type is consumed from different places. If the same variable was used on each arc from two different places, both places would need a token with the same value.

Variable declarations will be listed as they are encountered in the model. As a general guideline, most variables are named with the first letter of its colourset for non-lists, and the same letter plus the letter s for lists. On this submodule we have the following variables:

Listing 2.3: Client Application Variables

```

var msg: MESSAGE;
var msgs, msgs2: MESSAGES;
var es: EVENTS;

```

Queues

A lot of the places in this model rely on tokens being consumed in the same order they are produced, and to enforce this, queues are used, facilitated by list colour sets. To describe a list in SML, we use square brackets `[]`. By themselves they represent an empty list. To describe a populated list, we write each token inside the brackets separated by commas. An example of this is seen in Fig. 2.3 on the initial marking for the `Messages to be sent` place.

To process such a list, we use the `::` operator like this: `head::tail`, where `head` is the first element in the list, and `tail` is all the following elements. To concatenate two lists, we use the `^^` operator.

To emulate the queue behaviour we need, we use a list of a colour instead of using the actual colour we want in that place. When we want to take an element from the front of the queue, we use the `::` operator to bind the head and tail of the list to variables, and put only the tail back to the source place. When we want to append an element to a queue, we concatenate the queue using the `^^` operator with a new list containing only the new element. To improve readability of the model, these queue operations have one arc slightly dimmed, to emphasise the flow direction of data.

An example of this is seen in Fig. 2.3 on the arc from the `Send data` transition to the `Messages to be sent` place. On the left side, the list in the `Messages to be sent` place is bound to the two variables `msg` and `msgs`, and only `msgs` is put back. On the right side, we take the list in `Send Client Messages` and bind it to the variable `msgs2`, and send back a concatenation of this list and a new list containing `msg`, which was bound earlier in the incoming arc from `Messages to be sent`. The total effect is that we take the first element from `Messages to be sent`, and add it to the end of the list in `Send Client Message`.

Program Flow

The Request connection transition at the top is highlighted by a green border, which means that it is enabled and ready to fire. If we were to fire this transition, it would consume the `STRING` token, named `s`, produce a `MESSAGE` token containing `s`, and add it to the queue in the `Send Client Message` place.

Next, the Client Application waits for a `CONN_RESULT` token, and if this is equal to `success`, we add a `UNIT` token to the Active Connection place. If any other messages are received before this, they will wait at the `Receive Client Event` place, since the `Receive Data` transition will not fire unless the Active Connection place has a `UNIT` token.

If and when the Active Connection place has a `UNIT` token, the Client Application can start sending and receiving messages. A sample of messages has been predefined at the `Messages to be sent` place, but this would probably be dynamic in a real application.

Finally, if the Client Application receives a `MESSAGE` where the `OPERATOR` is `CLOSE`, nothing is put back into the Active Connection place, and the connection is effectively closed.

Kanskje dette kunne vært mer abstrakt? Argument for hvorfor modellert på denne måten.

2.3.3 The Client WebSocket Module

This module consists mostly of substitution transitions. The only processing being done directly on this module is masking of all websocket frames, which is required from the client by the protocol specification. The rest is plumbing between the submodules.

The new coloursets are for HTTP requests and responses, WebSocket frames, and packets. They have corresponding variables.

Hvordan funker mask funksjonen?

Listing 2.4: Library colour sets

```
colset HTTP_VERB = union GET + POST + PUT + DELETE + HEAD;
```



```

colset REQUEST_LINE = record
Verb: HTTP_VERB *
Path: STRING *
Version: STRING;

colset HEADER = product STRING*STRING;

colset HEADERS = list HEADER;

colset HTTPREQ = record
RequestLine: REQUEST_LINE *
Headers: HEADERS;

colset RESPONSE_LINE = record
Version: STRING *
Status: INT *
Message: STRING;

colset HTTPRES = record
ResponseLine: RESPONSE_LINE *
Headers: HEADERS;

colset MASK = list BYTE with 4..4;

colset MASKING = union
Nomask + Mask:MASK;

val opContinuation = 0x0;
val opText = 0x1;
val opBinary = 0x2
val opConnectionClose = 0x8;
val opPing = 0x9;
val opPong = 0xA;

colset WSFRAME = record
Fin: BIT *
Rsv1: BIT *
Rsv2: BIT *
Rsv3: BIT *
Opcode: INT *
Masked: BIT *
Payload_length: INT *
Masking_key: MASKING *
Payload: STRING;

colset WSFRAMES = list WSFRAME;

colset PACKET = union
HttpReq:HTTPREQ +

```

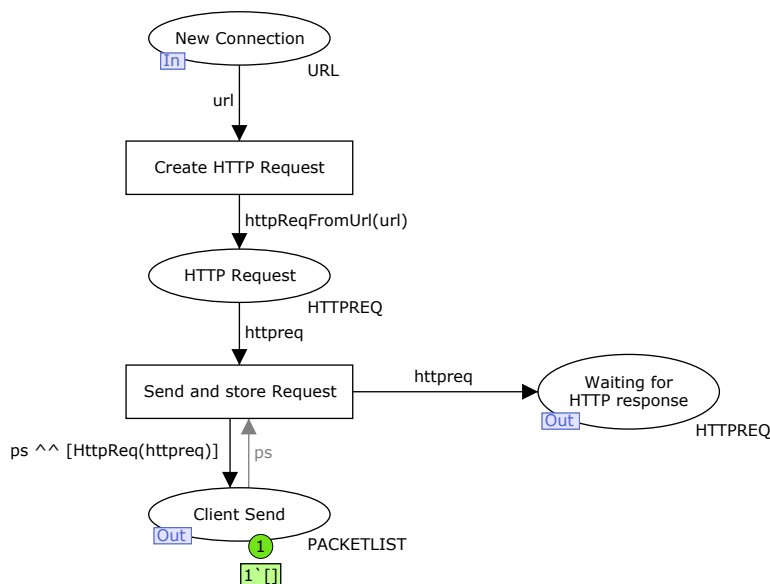


Figure 2.5: New Connection submodule

```

HttpRes:HTTPRES +
WsFrame:WSFRAME;

colset PACKETLIST = list PACKET;

```

Listing 2.5: Library variables

```

var wsframe: WSFRAME;
var wsframes, wsframes2: WSFRAMES;
var httpreq: HTTPREQ;
var httpres: HTTPRES;
var p: PACKET;
var ps, ps1, ps2: PACKETLIST;

```

There are also static values declared here, which correspond to WebSocket frame operation identifiers.

New connection

This module is fairly straightforward. We take the first `MESSAGE` in the list and create an HTTP request, which is then queued to be sent by the client, as well as keeping a copy of the request for validation purposes when the response arrives.

The transition here has a guard inscription enclosed in square brackets. It is used for two things: The first line limits the transition to only accepting messages with the `CONNECT` operation, while the second line binds the `url` variable by parsing the string in the message into a `URL` colour by using the `parseurl()` function. This function makes use of a utility function `split()` which splits a string on the first occurrence of a given character. To do this, it uses a recursive function `split2()` to iterate over the characters. The functions have to be declared in reverse order to satisfy dependency.

Listing 2.6: URL Declarations

```
colset URL = record
Protocol: STRING *
Host: STRING *
Port: INT *
Path: STRING;

var url: URL;

fun split2 (s, t, i) =
(* Recursively scan for character t in string s starting as
   position i, split
   if match *)
  let
    val ss = String.extract(s, i, NONE)
  in
    if String.isPrefix t ss then
      [substring(s, 0, i),
       String.extract(s, i + String.size t, NONE)]
    else split2(s, t, i+1)
  end

fun split (s, t) =
(* Split string s on character c *)
  if String.isPrefix t s then
    [String.extract(s, String.size t, NONE)]
  else if String.isSubstring t s then
    split2 (s, t, 1)
  else [s]

fun parseUrl (s) =
  let
    val proto'rest = split (s, "://")
    val proto'rest =
      if length proto'rest = 1
      then "ws" :: proto'rest
      else proto'rest
    val pro = List.hd(proto'rest)
```

```

val host'path = split (List.nth(proto'rest, 1), "/")
val pat = if length host'path = 2
  then "/" ^ List.nth(host'path, 1)
  else "/"

val host'port = split (List.hd(host'path), ":")
val hos = List.hd(host'port)

val port'default = case pro of
  "wss" => 443
  | _ => 80
val por = if length host'port = 1
  then port'default
  else let
    val port'str = List.nth(host'port, 1)
    val port'int'opt = Int.fromString port'str
  in
    Option.getOpt(port'int'opt, port'default)
  end
in
  {Protocol=pro,
   Host=hos,
   Port=por,
   Path=pat}
end;

```

The function `httpReqFromUrl` takes a URL argument and uses it to produce a `HTTPREQ` colour with headers according to the WebSocket protocol requirements. Not all headers defined by the protocol are used, because it is assumed this application is not browser based, and therefore specifying an origin does not make sense.

Listing 2.7: `httpReqFromUrl`

```

fun httpReqFromUrl (url:URL) =
{
  RequestLine={
    Verb=GET,
    Path=(#Path url),
    Version=httpVersion
  },
  Headers=[
    ("Host", (#Host url)),
    ("Upgrade", "websocket"),
    ("Connection", "Upgrade"),
    ("Sec-WebSocket-Key", (B64 nonce)),
    ("Sec-WebSocket-Version", "13")
  ]
}

```

Skriv hvorfor det er en rimelig antagelse. (Det er det kanskje ikke)

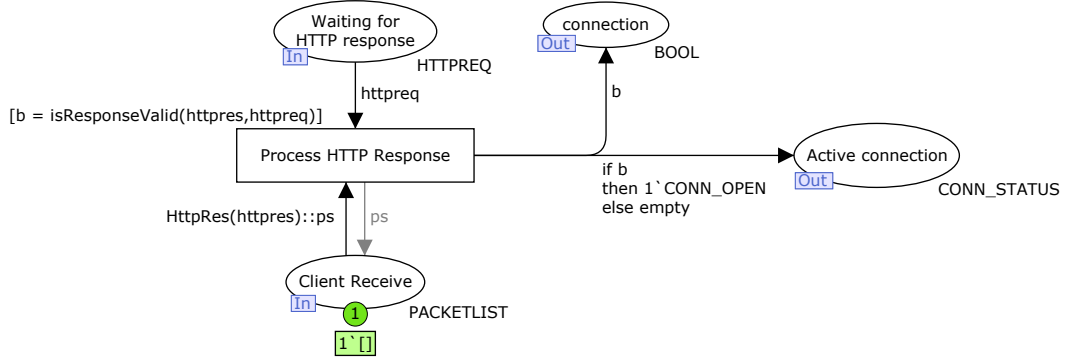


Figure 2.6: Process Response submodule

};

Process response

On this module, the transition fires when a HTTPRES colour is received from the server (and the HTTPREQ colour from earlier is still waiting in its place). The boolean variable *b* is bound to the result from the `isResponseValid` function, which checks if the server's reply is valid and conforms to the WebSocket protocol specification.

Listing 2.8: `isResponseValid`

```

fun isResponseValid (res:HTTPRES, req:HTTPREQ) = let
  val rline = #ResponseLine res
  val headers = #Headers res
  val accepttoken = generateAccept(
    getHeader("Sec-WebSocket-Key",
      (#Headers req)))
in
  #Status rline = 101 andalso
  getHeader("Upgrade", headers)
    = "websocket" andalso
  getHeader("Connection", headers)
    = "Upgrade" andalso
  getHeader("Sec-WebSocket-Accept", headers)
    = accepttoken
end

```

It sends the result directly to the Client App, and puts a token with colour `CONN_OPEN` in the Active Connection place if *b* is true.

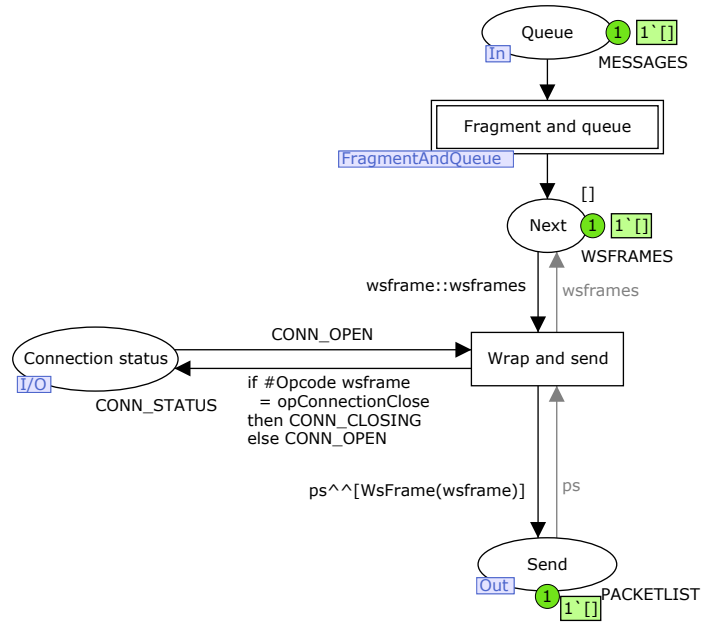


Figure 2.7: Wrap And Send submodule

Wrap and send

This module takes new messages, wraps and optionally fragments them in the Fragment and queue submodule (explained below), and sends them if there is an open connection. If a Close frame is being sent, the connection state will be changed to `CONN_CLOSING`, which prevents further sending.

Fragment and queue The top transition on this has two tasks: Making sure we do not process `CONNECT` messages, and filtering data and control frames into different places using the `isData` function.

Listing 2.9: `isData`

```

fun isData (msg:MESSAGE) =
  (#Op msg = TEXT) orelse
  (#Op msg = BINARY);
  
```

Control frames should never be fragmented, but data frames could be. This is taken care of by the `fragment` function, which also converts the message into a `WSFRAME` colour. The control frames are converted in the same way by the `wrapmsg` function. Both of these functions rely on the `wrap` function, and therefore have to be declared after it.

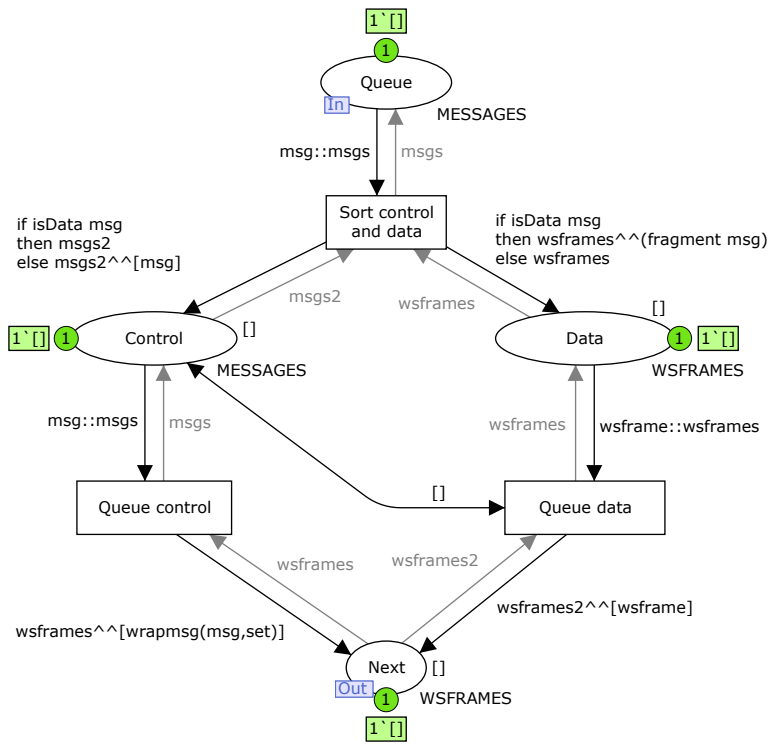


Figure 2.8: Fragment And Queue submodule

Listing 2.10: wrap wrapmsg and fragment

```

fun wrap (opc,payload,fin) = { Fin=fin,
    Rsv1=clear,
    Rsv2=clear,
    Rsv3=clear,
    Opcode=opc,
    Masked=clear,
    Payload_length=(String.size payload),
    Masking_key= Nomask,
    Payload=payload
}

fun wrapmsg (msg:MESSAGE,fin) =
    wrap(opSym2Hex(#Op msg),
        (#Message msg), fin);

fun fragment (msg:MESSAGE) = let
    fun loop (opc, s, acc) =
        if (String.size s) > fragSize
        then loop(
            opContinuation,
            String.extract(s,fragSize,NONE),
            acc^[wrap(opc,
                String.substring(s,0,fragSize),
                clear
            )]
        )
        else
            acc^[wrap(opc, s, set)];
    in
        loop(
            opSym2Hex(#Op msg),
            (#Message msg),
            [])
    end;

```

The WSFRAME colours are then queued randomly one by one from the data queue or the control frame queue. This allows control frames to be injected between the parts of a fragmented data frame, as required by the WebSocket protocol specification. Control frames are prioritised, by the prescence of a two-way arc from the Control place to the Queue data transition, inscribed with \square , which prevents it from being enabled if the list in the Control place is not empty. This priorisation is allowed but not required by the WebSocket protocol specification, but is included here to emphasise that control frames can be sent even between two fragmented data frames.

Unwrap and receive

Received WebSocket frames that arrive in the Packet Received place at the bottom can take three paths.

The first is to the left, and happens if the connection is in the `CONN_OPEN` state (checked on the arc) and the frame is not a close frame (checked in the guard of the Receive transition). The WebSocket frame is put in the Received WS Frame place, and if it is a Ping frame, a Pong frame is immediately queued for sending with identical message body.

The received frame is then checked on two points for fragmentation: If the Fin bit is set and the opcode is not continuation, it is not part of a fragmented message and converted directly to a `MESSAGE`. If either or both of those conditions are not true, this is part of a fragmented message and is processed in the Defrag submodule (explained below).

The second path a frame can take is to the right, and happens if the connection is in the `CONN_OPEN` state (checked on the arc) and the frame is a close frame (checked in the guard of the Receive transition). A close frame is created and set to be sent, and the connection state is changed to `CONN_CLOSED`, since we have both received and sent a close frame. Note that the packetlist from Client Send is not appended to but instead discarded, because we can not expect the other end to process any more frames other than a close frame since it has already sent a close frame of its own.

The third path is upwards and happens if the connection state is `CONN_CLOSING`, which means a close frame has been sent and we are waiting for a reply. Any payload is ignored, and the connection state stays the same until a close frame is received, in which case the connection state is set to `CONN_CLOSED`.

Defragmenting fragmented frames Frames that are part of a fragmented message should always arrive in a specific order, and correspond to each of the transitions on this module.

If the Fin bit is clear and the opcode is not continuation, this is the first frame in the series. A new `MESSAGE` is created with the opcode and payload from the WebSocket frame and put in the Buffer place.

If the Fin bit is clear and the opcode is continuation, this frame belongs in the middle of the sequence. The payload is appended to the `MESSAGE` in the Buffer place using the `append()` function.

Listing 2.11: `append`

```

fun append (msg:MESSAGE, s) =
  {Op = #Op msg,
   Message = (#Message msg)^s}

```

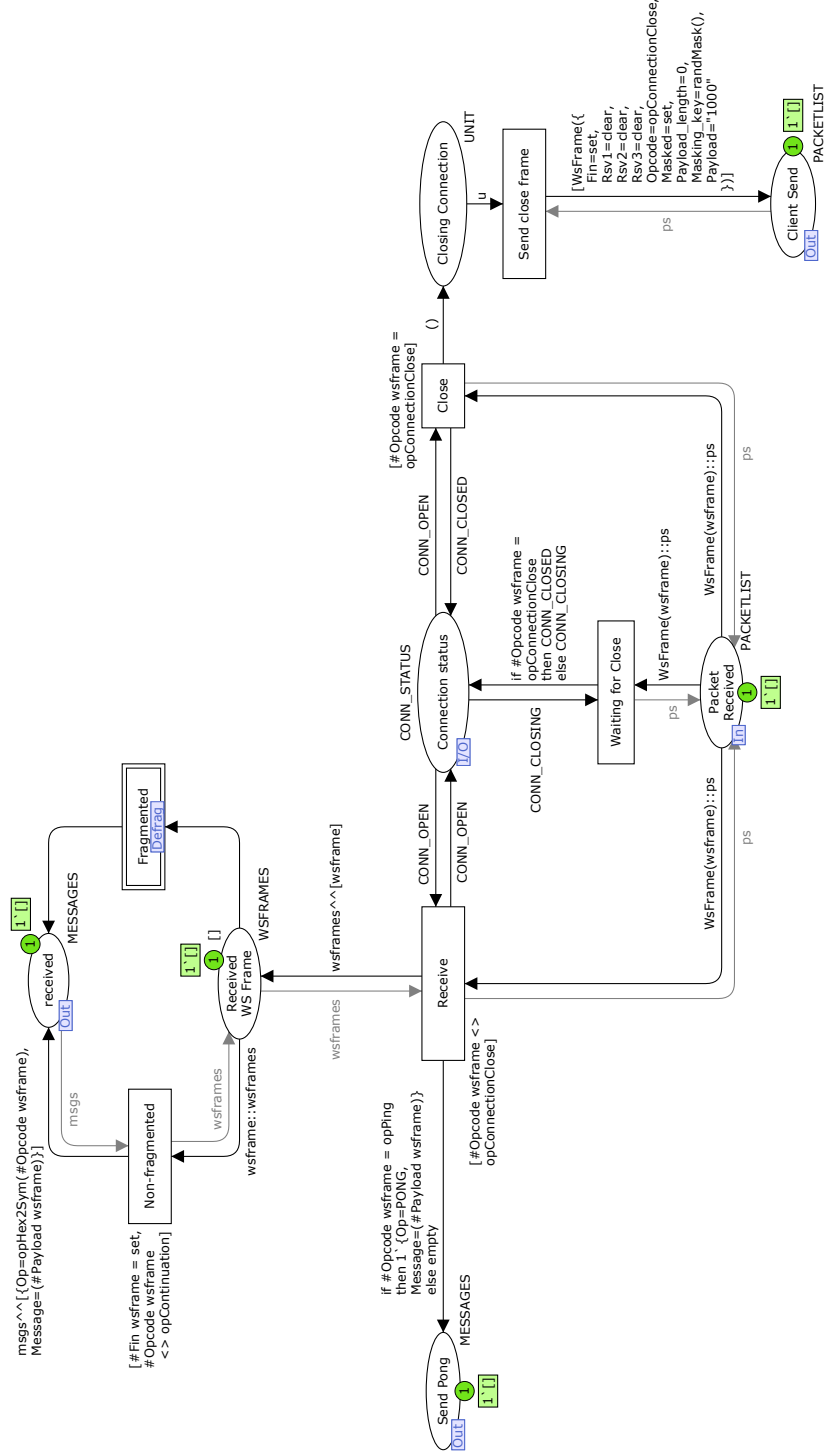


Figure 2.9: Unwrap And Receive submodule

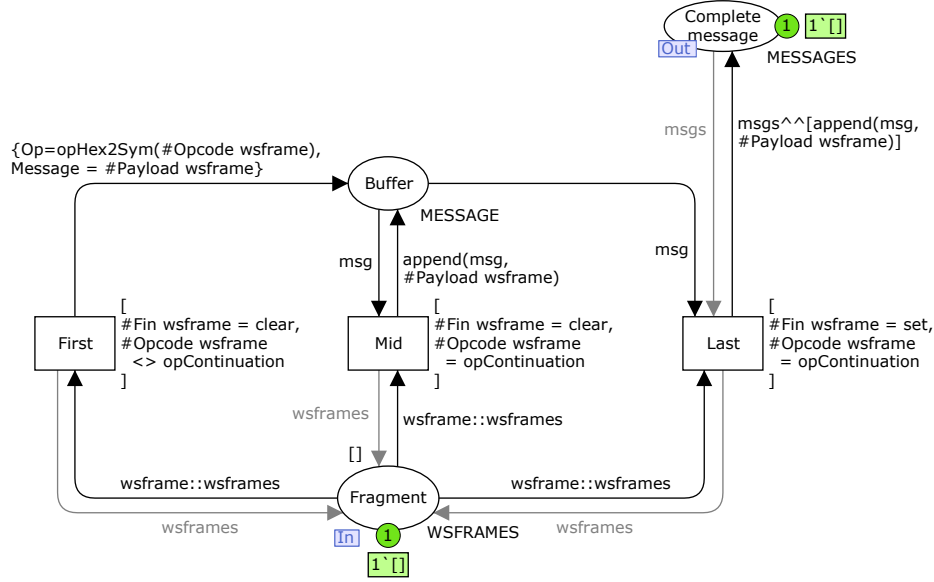


Figure 2.10: Defragment

If the Fin bit is set and the opcode is continuation, this is the last frame of the sequence. We append the payload to the message and put it in the final Completed message place.

Note that the WebSocket protocol does not allow fragmented messages to be interleaved, so it can be assumed that consecutive fragments belong to the same message. Fragment interleaving can be defined by subprotocols, but this is not relevant to this model.

2.3.4 Connection

Fig. 2.11 shows the connection; a very simple model. The packets that come from the Client Send place go to the Server Receive place, and from the Server Send place to the Client Receive place. The transportation of data between the client and the server as well as ensuring a stable connection is assumed to be taken care of by TCP, as the WebSocket protocol specifies, and is not necessary to model in detail to show how WebSocket works.

The packets are also not converted to pure bits or bytes. This decision was made to keep inspection easier during simulation.

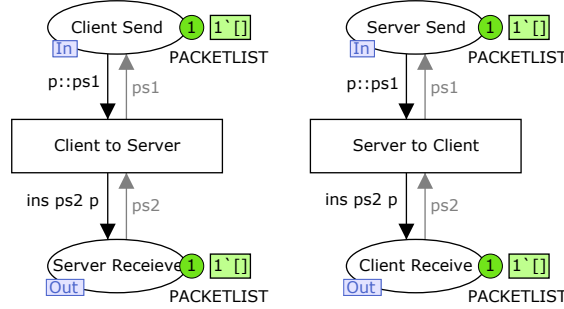


Figure 2.11: The Connection module

2.3.5 The Server WebSocket Layer

The Server WebSocket layer is very similar to the Client version. The main differences are that instead of masking outgoing frames, we are checking incoming frames for a mask and unmasking them, and that we are checking for incoming connections and replying to them based on what the Server Application decides.

The Server Application can send two types of coloursets: `MESSAGE` and `CONN_REPLY`. A special colourset has been made to accomodate this, `MSG_OR_CONN_REPLY`.

The submodules that belong to the `Wrap and Send` and the `Unwrap and Receive` substitution transitions are the same as the ones for the Client Library. To be more precise, the Client Library and Server Library both have instances of the same submodule, so that editing the submodule model affects both parent modules while simulating them lets them have different states.

Get Connection Request

This is a very abstract representation of how connection requests are received. The library simply sends a `ConnRequest` token to the app. In a real world app, much more info might have been sent, for example IP address and possibly authentication info, but for this model it is enough to show that some info is being sent, while the details are abstracted away.

Send Connection Response

When the Server Application has decided what to do with an incoming connection, it will send a `CONN_REPLY` to the library. If the answer is `accept`, we




Figure 2.12: The Server WebSocket Layer




Figure 2.13: Get Connection Request

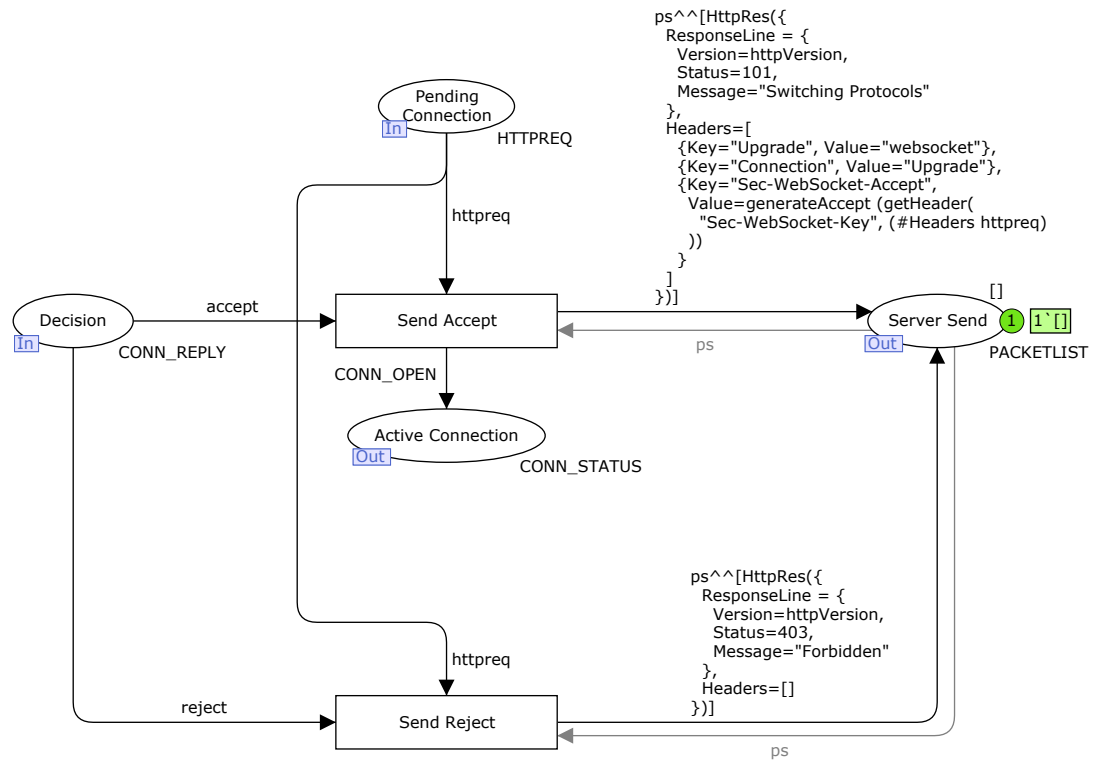


Figure 2.14: Send Connection Response

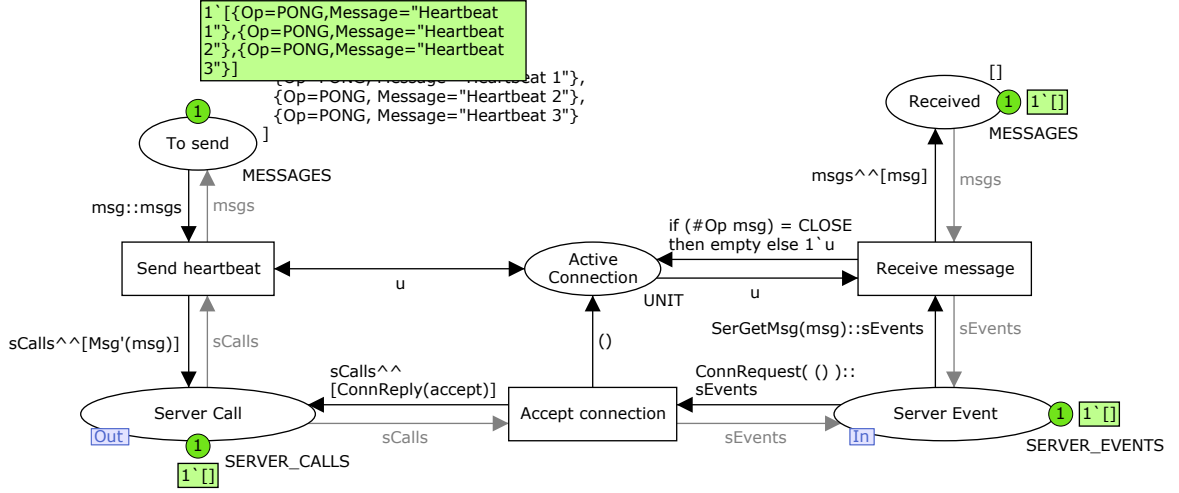


Figure 2.15: The Server Application Module

create a `CONN_OPEN` token in the Active Connection place and send a HTTP response back to the client, properly formatted according to the rules of the WebSocket Protocol. This involves generating a Sec-WebSocket-Accept header, which is done with the `generateAccept()` function:

Listing 2.12: generateAccept

```

fun B64 str =
  "B64("^str^")";

fun SHA1 str =
  "SHA1("^str^")";

fun generateAccept str =
  B64(SHA1(str^uuid));

```

The two functions `B64()` [Jos06] and `SHA1()` [iosat02] are abstract versions of the Base64 encoding algorithm and the SHA1 hashing algorithm. We felt it was unnecessary to actually implement these for the purpose of this model, and instead decided to simply wrap the string argument to show that it had been encoded or hashed.

2.3.6 Server Application

The Server Application has three tasks: Accept or reject incoming connections, and send and retrieve data. The colourset `MSG_OR_CONN_REPLY` that was

defined in Listing 2.1 is seen here. It has three queued messages, which illustrate the capability of PONG frames to be used as a heartbeat (without PING being involved). A real world application would have more logic here, but the interface to the library would be the same.

Chapter 3

State Space Analysis of the CPN Web Socket Protocol

One of the advantages of Coloured Petri Nets is the ability to conduct state space analysis, which can be used to obtain information about the behavioural properties of a CPN model, and which can be used to locate errors and increase confidence in the correctness of the CPN model.

3.1 State Spaces

A state space is a directed graph where each node represents a reachable marking (a state) and each arc represents an occurring binding element (a transition firing with values bound to the variables of the transition). CPN Tools by default generates the state space in breadth-first order.

TODO: Figur av SS initiell modell, med forklaring

Once generated, the state space can be visualised directly in CPN Tools. Starting with the node for the initial state, one can pick a node and show all nodes that are reachable from it, and in this way explore the state space manually. This can be very tedious and unmanageable for complex state spaces, though, and instead it is usually better to use queries to automate the analysis based on state spaces.

3.1.1 Strongly Connected Component graph

In graph theory, a strongly connected component (SCC) of a graph is a maximal subgraph where all nodes are reachable from each other. An SCC graph has a node for each SCC of the graph, connected by arcs determined

by the arcs in the underlying graph between nodes that belong to different SCCs. An SCC graph is acyclic, and an SCC is said to be trivial if it consists of only one node from the underlying graph.

By calculating the SCC graph of the state space, some of the further analysis becomes simpler and faster, such as determining reachability, cyclic behaviour, and checking so-called home and liveness properties.

3.1.2 Application of State Spaces

The biggest drawback of state space analysis is the size a state spaces may become very large. The number of nodes and arcs often grows exponentially with the size of the model configuration. This is also known as the state explosion problem.

This can be remedied by picking smaller configurations that encapsulate different parts of the system. This was necessary with the WebSocket Protocol model, as the complete state space took too long to generate.

Another aspect that must be considered prior to state space analysis is situations where tokens can be generated an unlimited amount of times on a place, thus making the state space infinite. This can be remedied by modifying the model to limit the number of simultaneous tokens in the offending place.

A model that incorporates random values is not always suited for computing a state space. The generated state space depends on the random values chosen, so the state space generator needs to be able to deterministically bind values to arc expression variables. Otherwise, a complete state space can not be achieved, since occurrence sequences might not converge after branching, and it's impossible to make sure all possible values have been considered.

For small color sets (generally defined as discrete sets usually with less than 100 possible values), binding random values to arc expressions can be done in two ways: By calling `ran()` on the colorset, or by using a free variable in the arc expression. The former, `ran()`, is non-deterministic and However, a free variable, which is a variable that does not get assigned a value in an expression, will also bind to a value picked at random from the color set during simulation, and also lets the state space generator pick each of the possible bindings from the values available in the colorset and thus generate all possible successive states.

Color sets that use values from a large or unbounded range, or from continuous ranges like floating point numbers, are considered large color sets, and using random values from such color sets can make it impossible

Diskutere mer hva som måtte gjøres med WebSocket, illustrasjon, konkret eksempel fra modellen

Vanskelig å formulere...

or impractical to generate a complete and correct state space. It can be worked around by providing single or tiny sets of fixed preset values. The CPN Tools manual has examples on how to do this.

For the WebSocket Protocol model, this was a problem for the masking key in WebSocket frames, which is supposed to be a random 4-byte string, giving 2^{32} or almost 4.3 billion possible values. To generate state spaces for this model, the randomisation function used was simply changed to always return four zeros.

3.2 State Space Report

Once a partial or complete state space has been generated, CPN Tools lets the user save a state space report as a textual document. The report is organised into parts that each describe different behavioural properties of the state space.

To explain each section of the state space report, a simple report for the WebSocket protocol has been generated, in a configuration where no messages are set to be sent. Thus, the only thing that will happen is that a connection will be established. Later in the chapter we will consider more elaborate configurations of the WebSocket protocol.

3.2.1 Statistics

The first section of the report describes general statistics about the state space.

State Space	
Nodes:	17
Arcs:	16
Secs:	0
Status:	Full
 Scc Graph	
Nodes:	17
Arcs:	16
Secs:	0

This state space has 17 possible markings, with 16 enabled transition occurrences connecting them. There is one more node than there are arcs, which means this graph is a tree.

The **Secs** field shows that it took less than one second to calculate this state space, while the **Status** field tells whether the report is generated from a partial or full state space. In this case the state space is fully generated.

We also see that the SCC Graph has the same number of nodes and arcs, meaning that there are no cycles in the state space (although this was already known from the fact that it is a tree).

3.2.2 Boundedness Properties

The second section describes the minimum and maximum number of tokens for each place in the model, as well as the actual tokens these places can have. The text has been reformatted and truncated (indicated by [...]) for readability.

Best Integer Bounds		
	Upper	Lower
ClientApplication		
Active_Connection	1	0
Conn_Result	1	0
Connection_failed	0	0
Messages_received	1	1
Messages_to_be_sent	0	0
[...]		

Many places show a lower and upper bound of 1. This shows a weakness in the approach of using lists to facilitate ordered processing of tokens: We cannot see the actual number of tokens that are in the place, because technically there is just a list there.

Best Upper Multi-set Bounds	
ClientApplication	
Active_Connection	1 '()
Conn_Result	1 'success
Connection_failed	empty
Messages_received	1 '[]
Messages_to_be_sent	empty
[...]	
ClientWebSocket	
Connection_status	1 'CONN_OPEN

```

[...]
ServerWebSocket
    Connection_Status
        1 'CONN_OPEN
[...]

Best Lower Multi-set Bounds
ClientApplication
    Active_Connection
        empty
    Conn_Result
        empty
    Connection_failed
        empty
    Messages_received
        1 '[]
    Messages_to_be_sent
        empty
[...]
ClientWebSocket
    Connection_status
        empty
[...]
ServerWebSocket
    Connection_Status
        empty
[...]

```

Apart from that, we see that both the client and the server has an open connection at some point, as the `Connection_status` place in the `ClientWebSocket` and `ServerWebSocket` modules have both had a `CONN_OPEN` token.

3.2.3 Home Properties

This section shows all home markings. A home marking is a marking that can always be reached no matter where we are in the state space.

Home Markings
[17]

Vise marking

We see that there is one such marking defined by node 17. From earlier we know that the state space is a tree, and if this node is always reachable it must be a leaf and all the other nodes must be in a chain. This tells us that there is only one possible sequence of transitions to establish a connection. We can then confidently say that the model works correctly with this configuration.

3.2.4 Liveness Properties

This section describes liveness of the state space. Some of the transitions have been omitted for readability.

```
Dead Markings
[17]
```

```
Dead Transition Instances
ClientApplication'Fail 1
ClientApplication'Receive_data 1
ClientApplication'Send_data 1
ClientWebSocket'Filter_messages 1
.....
```

```
Live Transition Instances
None
```

from where? for which?
formuler

A dead marking is a marking from where no other markings can be reached. In other words, there are no transitions for which there are enabled bindings, and the system is effectively stopped. For our example, we have a single dead marking, and it is the same as our home marking, confirming that this is a leaf node in the tree.

We also get a listing of dead transition instances, which are transitions that never have any enabled bindings in a reachable marking and are thus never fired. This can be useful to detect problems with a model, but in this example it is expected for many of the transitions, since we are not sending any kind of messages in the configuration considered.

Last, there are live transition instances. A transition is live if we from any reachable marking can find an occurrence sequence containing the transition. Our example has no such transition, which follows trivially from the fact that there is a dead marking.

The state space report also contains fairness properties, but this does not apply to our model since it contains no cycles. We will not explain more about this here, and instead refer to [JK09] chapter 7 for more information.

3.3 TODO:

Skrive om resten av analysene.

Flette inn dette:

3.3.1 Error discovery

An error in the model was discovered this way, in the Unwrap and Receive module, where the pong reply was adding a new list instead of appending to the old one in outgoing messages.

Flere detaljer - hvordan
viste fielen seg? Figur, før
og etter

Chapter 4

Technology and Foundations

One of the first decisions that had to be made for this thesis was whether to base the work on some existing platform or to create a new one from scratch. In this chapter we will describe the reasoning behind the design choices we made, and give an overview of the technologies that have been used.

TODO: Krav

4.1 Representing CPN Models

A (???) design decision is how to represent the CPN models. A simple but easy way of manipulating a CPN model is by representing it as a tree, with pages as nodes, and places, transitions and arcs as child nodes with properties describing how they connect in the actual CPN model. Simple tree editors are a feature of most GUI software platforms. Even so, we realised early that writing everything from scratch would take much longer than adapting an existing platform.

There are of course many complete implementations of Petri Net tools in different languages and toolkits, but few of them are open source, or written with extensibility in mind. If we were to base our work on an existing platform, it would have to be open and extendable.

To narrow our search, we limited our options to solutions in languages we had experience with: Java, c++/Qt and Ruby. Java is a popular language, and we already have Access/CPN, a part of the CPN Tools project, which can parse .cpn files and represent the model as Java objects.

The ePNK framework, an extendable framework for working with Petri Nets in a graphical manner, and that makes it possible to specify your own

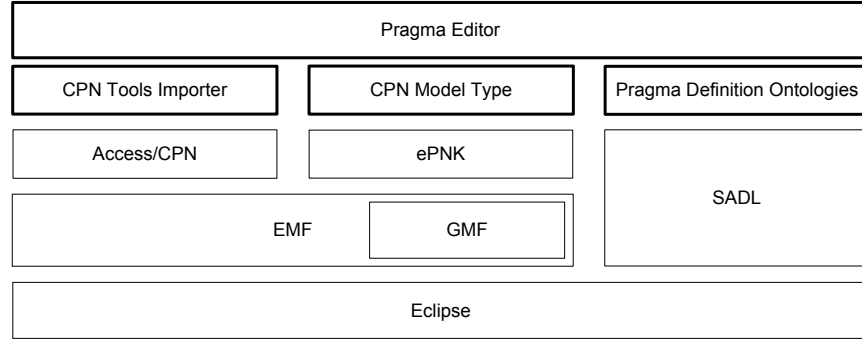


Figure 4.1: Application Overview Diagram

Petri Net type. It is built on the Eclipse Modeling Framework (EMF) (which Access/CPN is also built on).

We also needed a way to represent pragmatics. It was suggested to try an ontology-based approach, and we selected Semantic Application Design Language (SADL), another Eclipse plugin that lets us easily define and work with ontologies.

Fig. 4.1 shows the different elements that make up the application. The elements with bold frames are the ones newly created for this thesis, while the rest below are the existing solutions used and built upon. These will be described in the following sections, from the bottom up.

4.2 Eclipse IDE

Eclipse IDE [ecl] is an open source, cross-platform, polyglot development environment. Its plugin framework makes it highly extendable and customisable, and especially makes it easy for developers to quickly create anything from small custom macros, to advanced editors, to whole applications. The Eclipse IDE is open source, and part of the Eclipse Project, a community for incubating and developing open source projects.

The Eclipse IDE is built on the Eclipse Rich Client Platform (RCP) shown in Fig. 4.2. At the bottom of this we have the Platform Runtime, based on the OSGi framework, which provides the plugin architecture.

The Workspace
The Workbench




Figure 4.2: The Eclipse RCP

The Team plugin is a foundation for collaboration and versioning systems. It unifies many operations that are common between version control systems...

The Help plugin

Together these plugins form a basic generic IDE. Other plugins build on this to specialise the environment for a programming language and/or type of application.

The principal Eclipse distribution is the Eclipse Java IDE, which is one of the most popular tools for developing Java applications, from small desktop applications, to mobile apps for Android, to web applications, to enterprise-scale solutions.

Plugins are the building blocks of Eclipse, and there exists a wide range of plugins that add tools, functionality and services. For example, this thesis was written in \LaTeX using the Texlipse plugin, and managed with the Git version control system through the EGit plugin.

Publishing a custom plugin is simple. By packaging it and serving it on a regular web server, anyone can add the web server URL to the update manager in Eclipse, and it will let you download and install it directly, as well as enabling update notifications.

It is possible to package Eclipse with sets of plugins to form custom editions of Eclipse that are tailored for specific environments and programming languages. Aptana Studio is one example, aimed at Ruby on Rails and PHP development.

4.3 Eclipse Modeling Framework (EMF)

EMF is a framework for Model Driven Development (MDD) in Java. It is an Eclipse plugin that is part of the Eclipse Platform, and is open source. The principle of MDD is to define model structure by creating a metamodel to specify which entities can be created and how they relate to each other. By providing modeling and code generation tools, EMF allows developers to create model specifications (metamodels) that can be converted to Java classes, along with a set of adapter classes that enable viewing and command-based editing of the model, and a basic editor. These capabilities make EMF ideal for obtaining a simple editor that can be used to manipulate CPN models.

EKSEMPEL

4.3.1 Graphical Modeling Framework (GMF)

GMF builds on EMF to provide graphical viewing and editing of models. It uses metamodels created with EMF to generate implementations of views and editors that can create and edit the respective models. This can be used to create an editor that looks and works similarly to CPN Tools, and also lets you annotate it with pragmatics.

4.4 ePNK: Petri Net framework

ePNK is an Eclipse plugin both for working with standard Petri Net models, and a platform for creating new tools for specialised Petri Net types, which is exactly what we need for our annotated Petri Net type. It uses EMF and GMF to work with the Petri Net models and provide generic editors for custom Petri Net variants.

There are several reasons why ePNK is a good choice:

- It saves models using the ISO/IEC 15909 standard file format Petri Net Markup Language (PNML),
- It is currently actively developed,
- It is designed to be generic and easily extendable by creating new model types, and
- It includes both a tree editor and a graphical editor, provided through GMF.

referanse?

ePNK includes definitions for the core PNML model type, as well as two subtypes of Petri Nets. The first is P/T-Nets (Place/Transition Nets), which expand on the core model with a few key items: initial markings for places as integers, inscriptions on arcs, and constraining arcs to only go between a place and a transition (this is not enforced in PNML, as there are Petri Net variants that allow this).

The second type included with ePNK is High level Petri Nets (HLPNG). This type adds Structured Labels which are used to represent model declarations, initial markings, arc expressions and transition guards. These are parsed and validated using a syntax that is inspired from (but not the same as) CPN ML from CPN Tools. It is possible to write invalid data in these labels and still save the document, as they will only be marked as invalid by the editor to inform the user.

Neither of these two types conform exactly to the Coloured Petri Nets created by CPN Tools. HLPNG comes close, but is missing a few things like ports and sockets (RefPlaces can emulate this), and substituting transitions. Also, the structured labels are not compatible with CPN ML syntax from CPN Tools, and for our prototype, these structured labels are not necessary with regard to annotations. They might be useful in a future version, where for example pragmatics are available depending on things like the colorset of a place or the variables on an arc, but initially this is considered to be out of the scope of this thesis.

Our decision was therefore to develop our own Petri net type that matches the type supported by CPN Tools.

4.5 Access/CPN: Java interface for CPN Tools

CPN Tools has a sister project called Access/CPN. This is an EMF-based tool to parse .cpn files and represent them as an EMF-model. The .cpn files saved by CPN Tools are XML-based, which makes them easy to parse, but having an existing solution for this is preferable.

The model definition used by Access/CPN is very similar to that of ePNK.

TODO: Vurdere å utvide med egen algoritme som konverterer direkte til ePNK?

Diskutere mer i
implementation?

4.6 Semantic Application Design Language (SADL): Ontologies

Ontologies are a way to present information and meta-information so that it can be understood by computers. Essentially, this is done by defining classes that have properties, relations and constraints, and then present information with these classes.

There is a lot of ongoing research on this subject, especially to create a semantic web, that is extending web pages to provide meta-information about the content they contain and enabling software to understand it and reason about it. The Web Ontology Language (OWL) is the standard for representing ontologies.

Ref, sjekk om det er ISO

SADL is an Eclipse plugin that defines an english-like syntax for defining ontologies, and comes with a text editor that features syntax highlighting, parsing and validation. This is useful as we can possibly reuse the editor in our plugin for defining model-specific pragmatics. SADL ontologies can be compiled to OWL format.

It also has tools to parse and reason with these ontologies, which we will use to filter and validate which pragmatics are available for different model entities.

TODO: Lite eksempel på ontology.

4.7 Summary

After picking these technologies, since all the componets have Eclipse in common, it was an easy decision to develop our project as an Eclipse plugin. This also let us centralise all our development in Eclipse.

Chapter 5

Analysis and Design

5.1 Requirements

Load models created in CPN Tools Achieved with Access/CPN .cpn files are XML formatted

- Annotate model with pragmatics Validation

- Load sets of pragmatics to add to model

- Create model specific pragmatics on the fly

5.2 Test cases

(NYI, trenger eksempler) Simple protocol Kao-chow

5.3 Defining Pragmatics

5.4 The CPN model type for ePNK

5.5 Importing from CPN Tools

(Burde vært under Implementation?) Access/CPN uses many of the same class names as ePNK, making code difficult to read due to the need to write fully qualified classpaths. Does not parse graphics data. Large dependency for small feature. Consider rewriting importer to directly create ePNK objects

5.6 Creating annotations

Labels on nodes (and arcs?) and inscriptions and everything. Choose from list. Possibly write freehand with content assist. Validation, with problem markers (Eclipse feature).

5.7 Choosing Pragmatics Sets

Where to store? Model, Project, Plugin Model pros: Will need anyway for model-specific pragmatics Could be associated with net type as a property (like HLPNG does) or as a sub node somewhere Have URI string and version to check. cons: Keeping base pragmatics up to date a problem Namespace-based? Already required for ontology

5.7.1 Creating custom pragmatics

Dynamically supported in content assist If ontology-based, use SADL editor På sikt eget verktøy Hva kan det settes på, hvilke attributter har det. Oversette til ontologi

Bibliography

- [ecl] Eclipse ide.
- [FM11] I. Fette and A. Melnikov. The websocket protocol. Internet-Draft draft-ietf-hybi-thewebsocketprotocol-15, Internet Engineering Task Force, September 2011.
- [iosat02] National institute of standards and technology. FIPS 180-2, secure hash standard, federal information processing standard (FIPS), publication 180-2. Technical report, DEPARTMENT OF COMMERCE, August 2002.
- [JK09] Kurt Jensen and Lars M. Kristensen. *Coloured Petri Nets: Modelling and Validation of Concurrent Systems*. Springer, 2009.
- [Jos06] S. Josefsson. The Base16, Base32, and Base64 Data Encodings. RFC 4648 (Proposed Standard), October 2006.
- [Zim80] H. Zimmermann. OSI Reference Model—The ISO Model of Architecture for Open Systems Interconnection. *IEEE Transactions on Communications*, 28(4):425–432, 1980.

List of Figures

2.1	Sequence Diagram of the WebSocket protocol	4
2.2	Overview of CPN model of the WebSocket protocol	5
2.3	The Client Application	9
2.4	The Client WebSocket Module	12
2.5	New Connection submodule	14
2.6	Process Response submodule	17
2.7	Wrap And Send submodule	18
2.8	Fragment And Queue submodule	19
2.9	Unwrap And Receive submodule	22
2.10	Defragment	23
2.11	The Connection module	24
2.12	The Server WebSocket Layer	25
2.13	Get Connection Request	26
2.14	Send Connection Response	26
2.15	The Server Application Module	27
4.1	Application Overview Diagram	38
4.2	The Eclipse RCP	39

Listings

2.1	Overview colour sets	7
2.2	Simple Colourset Variables	8
2.3	Client Application Variables	10
2.4	Library colour sets	11
2.5	Library variables	14
2.6	URL Declarations	15
2.7	httpReqFromUrl	16
2.8	isResponseValid	17
2.9	isData	18
2.10	wrap wrapmsg and fragment	20
2.11	append	21
2.12	generateAccept	27